

FreePOPs Manual

5th July 2004

Contents

1	Introduction	3
1.1	Usage situations	3
1.2	Features	3
1.3	Plugins	4
2	History	4
3	FreePOPs configuration file	5
4	FreePOPs command line parameters	5
5	Email client configuration	6
6	Plugins	6
6.1	libero.lua	6
6.2	tin.lua	7
6.3	lycos.lua	7
6.4	popforward.lua	7
6.5	aggregator.lua	7
6.6	flatnuke.lua	8
6.7	kernel.lua	9
7	Creating a plugin	9
7.1	Plugins overview	9
7.1.1	The interface between the C core and a plugin	10
7.1.2	The interface between a plugin and the C core	11
7.2	The art of writing a plugin (plugins tutorial)	12
7.2.1	(step 1) the skeleton	12
7.2.2	(step 2) the login	13
7.2.3	(step 3) getting the list of messages	18

7.2.4 (step 4) the common functions	22
7.2.5 (step 5) deleting messages	23
7.2.6 (step 6) downloading messages	24
7.2.7 (step 7) test it	26
7.2.8 (step 8) the so mentioned last part of the tutorial	26
8 Submitting a bug	31
9 Authors	32
9.1 Developers	32
10 Thanks	32

1 Introduction

FreePOPs is a POP3 daemon plus a LUA interpreter and some extra libraries for HTTP and HTML parsing. Its main purpose is translating local POP3 requests to remote HTTP actions on the supported web-mails, but it is really more flexible: for example there is a plugin to read news from a website as if they were mails in a mailbox. You can easily extend FreePOPs on the fly, without even restarting it; you can add a plugin or modify an existing one simply changing the script file since the plugins are written in LUA and are interpreted on the fly.

1.1 Usage situations

FreePOPs can be useful in some situations, here we give the most obvious ones:

- You are behind a firewall that closes the 110 port but you need to read your mail and the web-mail of your mail provider sucks.
- Your mail provider does not allow you to access your mailbox with POP3 protocol, but only through the web-mail service.
- You prefer looking at your mailbox instead of browsing some websites news.
- You have to develop a POP3 server in less than a week and you want it to be reasonably fast and not so memory consuming.
- You are not a C hacker, but you want to benefit from a fast POP3 server frontend written in C and you want not to waste a month in writing the backend in C. LUA is a really simple and tiny language, one week is enough to learn it in a way that allows you to use it productively.

1.2 Features

FreePOPs is the only software I know with these features:

- POP3 server RFC compliant (not full featured but compliant).
- Portable (written in C and LUA that is written in C, so everything is written in the most portable language around the world).
- Small (in the sense of resources usage) and reasonably fast.
- Extremely extensible on the fly using a simple and powerful language.
- Pretty documented.
- Released under the GNU/GPL license (this means FreePOPs is Free Software).

1.3 Plugins

These are the plugins currently included in FreePOPs:

libero.lua This plugin fully supports the `www.libero.it` webmail for mailboxes with domains like `@libero.it`, `@iol.it`, `@inwind.it`, `@blu.it`.

tin.lua This plugin fully supports the `communicator.virgilio.it` webmail for mailboxes with domains like `@tin.it`, `@virgilio.it`.

lycos.lua This plugin supports the `mail.lycos.it` webmail for mailboxes with `@lycos.it` domains.

popforward.lua This is a plugin mainly used for testing some FreePOPs modules. It acts as a POP3 forward, it simply works as a man in the middle of you and a real POP3 server. This allowed us to test FreePOPs without having any real plugin already written. You may consider using it to masquerade a really buggy POP3 server that can be easily compromised through malicious requests. Obviously we suggest you to examine properly this plugin, and hack on it to prevent malicious requests to your server.

aggregator.lua Many sites provide an RSS backend for indexing their news. This plugin makes this RSS behave as a mailbox in which you can find a mail for each news.

flatnuke.lua This is a more powerful aggregator for FlatNuke-based¹ websites, and allows the download of the whole news body.

kernel.lua This is a plugin to monitor the news about the Linux kernel through ChangeLogs.

2 History

FreePOPs was not born from scratch. A similar project (only in the main usage situation) is LiberoPOPs.

The ancestor of FreePOPs is completely written in C for some uninteresting reasons. LiberoPOPs supports “plugins” but in a more static and complex way. The POP3 server frontend could be attached to a backend written in C, this means you have to recompile and restart LiberoPOPs each time to change a line in a plugin. Another interesting point is that LiberoPOPs was created from scratch in a really short time (you have to be Italian and use a `@libero.it` mail address to understand why), this means it was born with a lot of bugs and FIX-ME in the code.

¹See [HTTP://flatnuke.sourceforge.net](http://flatnuke.sourceforge.net) for the project homepage

The LiberoPOPs project had a quick success, because everybody needed it, and this means we had a lot of users. In the opensource (and also Linux) philosophy you have to release frequently and this was exactly what we did: we used to release every two days. We were working not with Unix users, nor hackers, but mostly with Win32 users. Suddenly we realized that they were lazy/bored of updating the software every 2 days. The ugly Win-world has taught them that software auto-updates, auto-install and even auto-codes probably.

We tried to solve this pulling out of the C engine most of the change-prone code, but this was really hard since C is not thought to do this. After LiberoPOPs had stabilized we started to think how to solve this.

A scripting/interpreted embedded language seemed to me a nice choice and after a long search on the net and on the newsgroup of my university I found LUA.. This is not the place for telling the world how good this small language is and I won't talk more about it here. Integrating the LUA interpreter in LiberoPOPs was not so hard and FreePOPs is the result. Now it is really easier to write/test a plugin and (even if it is not implemented yet) an auto-update facility is really easy to code since there is no need to recompile the C core in most cases.

3 FreePOPs configuration file

FreePOPs doesn't really need a configuration. Most users shouldn't change the configuration file. In case you are a developer or a really curious user the configuration file is `config.lua`, placed in the program directory under win32 or in `/etc/freepops/` in a posix environment.

Later you will learn that plugins are associated with a mail address domain, and some of these plugins are aliased to other domains to make it easier to fetch some news from some sites. Read the plugin documentation for more info about them, and maybe send as a mail with your new alias if you want it to be integrated in the next FreePOPs release.

Since version 0.0.11 the `config.lua` file has a policy section. In this section you can ban or allow classes of mail addresses. This may be useful to network administrators.

4 FreePOPs command line parameters

The real FreePOPs configuration is made trough command line arguments. They are described in depth in the man page in Unix environments, here we present only the most useful:

- p <port>, -port <port>** By default FreePOPs binds on port 2000. To alter this behaviour just use this switch.
- P <host>:<port>, -proxy <host>:<port>** To tell FreePOPs which is your HTTP proxy.
- v, -verbose, -w, -veryverbose** This tells FreePOPs to log some interesting info for bug reporting.
- t <num>, -threads <num>** FreePOPs is able to manage multiple connections, up to num. Default is 5.

In posix environments like Debian GNU/Linux you can start FreePOPs at boot time as a standard service. In this case the command line switches are stored in `/etc/default/freepops`, in some rpm based systems you should find the same file as `/etc/sysconfig/freepops`.

5 Email client configuration

To configure your email client you must change the pop3 server settings. Usually you must use localhost as the pop3 host name, and 2000 as the pop3 port. In case you install FreePOPs in another computer of your LAN, you should use the host's name instead of localhost, while in case you changed the default port with the `-p` switch you will have to use that same port in your email client. You always have to use a full email address as username, for example `something@libero.it` instead of only `something`. This is because FreePOPs chooses the plugin to load by looking at your username. Later we will present all the plugins and their associated domains.

6 Plugins

Here we give a detailed description of each plugin.

6.1 **libero.lua**

This plugin allows you to read the mails you have in a `@libero.it`, `@iol.it`, `@inwind.it` and `@blu.it` mailbox. This means you can still use your favorite mail reader instead of using the webmail. This plugin acts as a browser that surfs your webmail account and make it appear as a POP3 server. For more info on this plugin you may read the LiberoPOPs (FreePOPs ancestor) website at <http://liberopops.sourceforge.net>

To use this plugin you have to use your full email address as the username and your real password as the password.

6.2 **tin.lua**

This is the webmail support for @virgilio.it and @tin.it mailboxes. To use this plugin you have to use your full email address as the username and your real password as the password.

6.3 **lycos.lua**

This is the webmail support for @lycos.it mailboxes. To use this plugin you have to use your full email address as the username and your real password as the password.

6.4 **popforward.lua**

This plugin was developed to test FreePOPs before any real plugins were written. It simply forwards to a real POP3 server your local requests. This can be used for masquerading a buggy POP3 server, but if you think you need this you should examine carefully the plugin code and add checks to improve the detection/avoidance of malicious requests, since the plugin was not born with security in mind.

To use this plugin you have to edit the config.lua file. This because we can't add all the existing POP3 server to this file :) The plugin wants two arguments, the POP3 host and the port (usually 110) on which the server listens. This is an example of a configuration line for this plugin, in which all email addresses of the @virgilio.it domain are forwarded to in.virgilio.it:110:

```
-- popforward plugin
freepops.MODULES_MAP["virgilio.it"] = {
    name="popforward.lua",
    args={
        port=110,
        host="in.virgilio.it"
    }
}
```

6.5 **aggregator.lua**

Usually you can benefit from the W3C's RSS format when you read some website news. The RSS file indexes the news, providing a link to them. This plugin can make your mail client see the RSS file as a mailbox from which you can download each news as if it was a mail message. The only limitation is that this plugin can fetch only a news summary plus the news link.

To use this plugin you have to use a casual username with the @aggregator suffix (ex: foo@aggregator) and as the password the URL of the RSS file(ex:

<http://www.securityfocus.com/rss/vulnerabilities.xml>). For your commodity we added some alias for you. This means you have not to search by hand the URL of the RSS file. We added some domain, for example `@securityfocus.com`, that can be used to directly use the aggregator plugin with these website. To use these alias you have to use a username in the form `something@aggregatordomain` and a casual password. This is the list of aliases for the aggregator plugin.

aggregatordomain	description
freepops.rss.en	http://freepops.sourceforge.net/ news (English)
freepops.rss.it	http://freepops.sourceforge.net/ news (Italian)
flatnuke.sf.net	http://flatnuke.sourceforge.net/ news (Italian)
ziobudda.net	http://ziobudda.net/ news (both Italian and English)
punto-informatico.it	http://punto-informatico.it/ news (Italian)
linuxdevices.com	http://linuxdevices.com/ news (English)
gaim.sf.net	http://gaim.sourceforge.net/ news (English)
securityfocus.com	http://www.securityfocus.com/ new vulnerabilities (English)
games.gamespot.com	http://www.gamespot.com/ computer games news (English)
news.gamespot.com	http://www.gamespot.com/ GameSpot news (English)

6.6 flatnuke.lua

This plugin is an aggregator plugin specialized in the websites made with the FlatNuke² content management system, or other sites that use the same news format like the FreePOPs website. Since in a FlatNuke site the news are stored in plain xml files this plugin is able to fetch the whole news, and not only the headings as the aggregator plugin does. This is really useful if you don't want to surf the website to read the news.

To use this plugin you have to use a username with the `@flatnuke` domain (ex: `something@flatnuke`) and a flatnuke homepage URL as the password (ex: <http://flatnuke.sourceforge.net/>, no need for the RSS file URL since FlatNuke puts the RSS in a fixed and well known position. There are some alias for FlatNuke sites, see the aggregator plugin documentation to know what this means):

aggregatordomain	description
freepops.en	http://freepops.sourceforge.net/ full news (English)
freepops.it	http://freepops.sourceforge.net/ full news (Italian)
flatnuke.it	http://flatnuke.sourceforge.net/ full news (Italian)

²[HTTP://flatnuke.sourceforge.net](http://flatnuke.sourceforge.net)

6.7 **kernel.lua**

This plugin helps in staying up to date with the Linux kernel releases. <http://kernel.org> is the official page with the Linux kernel releases, each with its ChangeLog. You should use something@kernel.org to receive news about every tree, something@kernel.org or something@kernel.org.26 for a specific tree. Password is not used, type a random string.

7 Creating a plugin

Two sections follow, the first is a quick overview of what a plugin has to do, the latter is a more detailed tutorial. Before proceeding I suggest you to read some stuff that is at the base of plugin writing:

1. Since plugins are written in LUA you must read at least the lua tutorial ([HTTP://lua-users.org/wiki/LuaTutorial](http://lua-users.org/wiki/LuaTutorial)), many thanks to the guys who wrote it. LUA is a quite simple scripting language, easy to learn, and easy to read. If you are interested in this language you should read THE book about LUA ("Programming in Lua" by Roberto Ierusalimsky [HTTP://www.inf.puc-rio.br/~roberto/book/](http://www.inf.puc-rio.br/~roberto/book/)). It is a really good book, believe me.
2. Since we have to implement a POP3 backend you should know what POP3 is. The rfc is number 1939 and is included in the doc/ directory of the source package of FreePOPs, but you can fetch it from the net [HTTP://www.ietf.org/rfc/rfc1939.txt](http://www.ietf.org/rfc/rfc1939.txt).
3. Read carefully this tutorial, it is fairly a good tutorial, but is better than nothing.
4. The website contains, in the doc section, a quite good documentation of the sources. You should keep a web browser at the LUA modules documentation page while writing a plugin.
5. After creating a prototype, you should read a full featured plugin. The `libero.lua` plugin is really well commented, start from it.
6. Remember that this software has an official forum ([HTTP://freepops.diludovico.it](http://freepops.diludovico.it)) and some authors you may ask for help.

7.1 Plugins overview

A plugin is essentially a backend for a POP3 server. The plugins are written in LUA³ while the POP3 server is written in C. Here we examine the interfaces

³The language website is [HTTP://www.lua.org](http://www.lua.org)

between The C core and the LUA plugins.

7.1.1 The interface between the C core and a plugin

As we explained before the C POP3 frontend has to be attached to a LUA backend. The interface is really simple if you know the POP3 protocol. Here we only summarize the meaning, but the RFC 1939 (included in the `doc/` directory of the source distribution) is really short and easy to read. As your intuition should suggest the POP3 client may ask the pop3 server to know something about the mail that is in the mailbox and eventually retrieve/delete a message. And this is exactly what it does.

The backend must implement all the POP3 commands (like USER, PASS, RETR, DELE, QUIT, LIST, ...) and must give back to the frontend the result. Let us give a simple example of a POP3 session taken from the RFC:

```
1  S: <wait for connection on TCP port 110>
2  C: <open connection>
3  S:  +OK POP3 server
4  C:  USER linux@kernel.org
5  S:  +OK now insert the password
6  C:  PASS gpl
7  S:  +OK linux's maildrop has 2 messages (320 octets)
8  C:  STAT
9  S:  +OK 1 320
10 C:  LIST
11 S:  +OK 2 messages (320 octets)
12 S:  1 320
13 S:  .
14 C:  RETR 1
15 S:  +OK 120 octets
16 S:  <the POP3 server sends message 1>
17 S:  .
18 C:  DELE 1
19 S:  +OK message 1 deleted
20 C:  QUIT
21 S:  +OK dewey POP3 server signing off (maildrop empty)
22 C:  <close connection>
23 S:  <wait for next connection>
```

In this session the backend will be called for lines 4, 6, 8, 10, 14, 18, 20 (all the C: lines) and respectively the functions implementing the POP3 commands will be called this way

```
user(p, "linux@kernel.org")
pass(p, "gpl")
```

```

stat(p)
list_all(p)
retr(p,1)
dele(p,1)
quit_update(p)

```

Later I will make clear what `p` is. I hope we'll remove it making it implicit for complete transparency. It is easy to understand that there is a 1-1 mapping between POP3 commands and plugin function calls. You can view a plugin as the implementation of the POP3 interface.

7.1.2 The interface between a plugin and the C core

Let us take in exam the call to `pass(p, "linux@kernel.org")`. Here the plugin should authenticate the user (if there is a sort of authentication) and inform the C core of the result. To achieve this each plugin function must return an error flag, to be more accurate one of these errors:

Code	Meaning
POPSERVER_ERR_OK	No error
POPSERVER_ERR_NETWORK	An error concerning the network
POPSERVER_ERR_AUTH	Authorization failed
POPSERVER_ERR_INTERNAL	Internal error, please report the bug
POPSERVER_ERR_NOMSG	The message number is out of range
POPSERVER_ERR_LOCKED	Mailbox is locked by another session
POPSERVER_ERR_EOF	End of transmission, used in the <code>popserver_callback</code>
POPSERVER_ERR_TOOFAST	You are not allowed to reconnect to the server now, wait a bit and retry
POPSERVER_ERR_UNKNOWN	No idea of what error I've encountered

In our case the most appropriate error codes are `POPSERVER_ERR_AUTH` and `POPSERVER_ERR_OK`. This is a simple case, in which an error code is enough. Now we analyze the more complex case of the call to `list_all(p)`. Here we have to return an error code as before, but we have also to inform the C core of the size of all messages in the mailbox. Here we need the `p` parameter passed to each plugin function (note that parameter may become implicit in the future). `p` stands for the data structure that the C expects us to fill calling appropriate functions like `set_mailmessage_size(p, num, size)` where `num` is the message number and `size` is the size in bytes. Usually it is really common to put more functions all together. For example when you get the message list page in a webmail you know the number of the messages, their size and uidl so you can fill the `p` data structure with all the informations for LIST, STAT, UIDL.

The last case that we examine is `retr(p,num,data)`. Since a mail message can be really big, there is no pretty way of downloading the entire message without making the mail client complain about the server death. The solution is to use a callback. Whenever the plugin has some data to send to the client he should call the `popserver_callback(buffer,data)`. `data` is an opaque structure the popserver needs to accomplish its work (note that this parameter may be removed for simplicity). In some cases, for example if you know the message is small or you are working on a fast network, you can fetch the whole message and send it, but remember that this is more memory consuming.

7.2 The art of writing a plugin (plugins tutorial)

In this section we will write a plugin step by step examining each important detail. We will not write a real and complete plugin since it may be a bit hard to follow but we will create an ad-hoc webmail for our purposes.

7.2.1 (step 1) the skeleton

The first thing we will do is copy the `skeleton.lua` file to `foo.lua` (since we will write the plugin for the `foo.xx` webmail, `xx` stands for a real domain, but I don't want to mention any websites here...). Now with your best editor (I suggest vim under Unix and scintilla for win32, since they support syntax highlights for LUA, but any other text editor is OK) open `foo.lua` and change the first few lines adding the plugin name, version, your name, your email and a short comment in the proper places.

```
-- ***** --
-- FreePOPs @--put here domain-- webmail interface
--
-- $Id: manual.lyx,v 1.28 2004/07/05 16:39:30 gareuselesinge Exp $
--
-- Released under the GNU/GPL license
-- Written by --put Name here-- <--put email here-->
-- ***** --

PLUGIN_VERSION = "--put version here--"
PLUGIN_NAME = "--put name here--"
```

Now we have an empty plugin, but it is not enough to start hacking on it. We need to open the `config.lua` file (in the win32 distribution it is placed in the main directory, while in the Unix distribution it is in `/etc/freepops/`; other copies of this file may be included in the distributions, but they are backup copies) and add a line like this

```
-- foo plugin
freepops.MODULES_MAP["foo.xx"] = {name="foo.lua"}
```

at the beginning of the file. Before ending the first step you should try if the plugin is correctly activated by FreePOPs when needed. To do this we have to add few lines to `foo.lua`, in particular we have to add an error return value to `user()`.

```
-- -----
-- Must save the mailbox name
function user(pstate,username)
    return POPSERVER_ERR_AUTH
end
```

Now the `user` function always fails, returning an authentication error. Now you have to start FreePOPs (if it is already running you don't have to restart it) and start telnet (under win32 you should open a DOS prompt, under Unix you have the shell) and type `telnet localhost 2000` and then type `user test@foo.xx`.

```
tassi@garfield:~$ telnet localhost 2000
Trying 127.0.0.1...
Connected to garfield.
Escape character is '^]'.
+OK FreePOPs/0.0.10 pop3 server ready
user test@foo.xx
-ERR AUTH FAILED
Connection closed by foreign host.
```

The server respond closing the connection and printing an authorization failed message (thats OK, since the `user()` function of our plugin returns this error). In the standard error file (the console under Unix, the file `stderr.txt` under Windows) are printed error messages, don't care about them now.

7.2.2 (step 2) the login

The login procedure is the first thing we have to do. The POP3 protocol has 2 commands for login, *user* and *pass*. First the client does a *user*, then it tells the server the password. As we have already seen in the overview this means that first `user()` and the `pass()` will be called. This is a sample login:

```
tassi@garfield:~$ telnet localhost 2000
Trying 127.0.0.1...
Connected to garfield.
```

```

Escape character is '^]'.
+OK FreePOPs/0.0.10 pop3 server ready
user test@foo.xx
+OK PLEASE ENTER PASSWORD
pass hello
-ERR AUTH FAILED

```

If you start FreePOPs with the `-w` switch you should read this on standard error/standard output

```

freepops started with loglevel 2 on a little endian machine.
Cannot create pid file "/var/run/freepopdsd.pid"
DBG(popserver.c, 162): [5118] ?? Ip address 0.0.0.0 real port 2000
DBG(popserver.c, 162): [5118] ?? Ip address 127.0.0.1 real port 2000
DBG(popserver.c, 162): [5118] -> +OK FreePOPs/0.0.10 pop3 server ready
DBG(popserver.c, 162): [5118] <- user test@foo.xx
DBG(log_lua.c, 83): (@src/lua/foo.lua, 37) : FreePOPs plugin 'Foo web mail' version '0.
*** the user wants to login as 'test@foo.xx'
DBG(popserver.c, 162): [5118] -> +OK PLEASE ENTER PASSWORD
DBG(popserver.c, 157): [5118] <- PASS *****
*** the user inserted 'hello' as the password for 'test@foo.xx'
DBG(popserver.c, 162): [5118] -> -ERR AUTH FAILED
AUTH FAILED
DBG(threads.c, 81): thread 0 will die

```

and the plugin has been changed a bit to store the user login and print some debug info. This is the plugin that gave this output:

```

foo_globals= {
    username="nothing",
    password="nothing"
}
-----
-- Must save the mailbox name
function user(pstate,username)
    foo_globals.username = username
    print("*** the user wants to login as '"..username.."'")
    return POPSERVER_ERR_OK
end
-----
-- Must login
function pass(pstate,password)
    foo_globals.password = password
    print("*** the user inserted '"..password..
        "' as the password for '"..foo_globals.username.."'")

```

```

        return POPSERVER_ERR_AUTH end
-----
-- Must quit without updating
function quit(pstate)
    return POPSERVER_ERR_OK
end

```

Here we have some important news. First the `foo_globals` table that will contain all the globals (values that should be available to successive function calls) we need. So far we have the username and password there. The `user()` function now stores the passed username in the `foo_globals` table and prints something on standard output. The `pass()` function now stores the password in the global table and prints some stuff. The `quit()` simply returns `POPSERVER_ERR_OK` to make FreePOPs happy.

Now that we know how FreePOPs will act during the login we have to implement the login in the webmail, but first uncomment the few lines in the `init()` function (that is called when the plugin is started) that loads the `browser.lua` module (the module we will use to login in the webmail). Here is the webmail login page viewed with Mozilla and the source of the page (with Mozilla it is Ctrl-U).



```

<html>
<head>
<title>foo.xx webmail login</title>
</head>
<body style="background-color : grey; color : white">
<h1>Webmail login</h1>
<form name="webmail" method="post" action="http://localhost:3000/">
login: <input type="text" size="10" name="username"> <br>
password: <input type="password" size="10" name="password"> <br>
<input type="submit" value="login">
</form>
</body>
</html>

```

Here we have 2 input fields, one called username and one called password. When the user clicks login the web browser will POST to `HTTP://localhost:3000/` the form contents (I used a local address because of commodity, but it should be something like `HTTP://webmail.foo.xx/login.php`). This is what the browser sends:

```
POST / HTTP/1.1
Host: localhost:3000
User-Agent: Mozilla/5.0 (X11; U; Linux i686; en-US; rv:1.6) Gecko/20040614 Firefox/0.8
Accept-Language: en-us,en;q=0.5
Accept-Encoding: gzip,deflate
Accept-Charset: ISO-8859-1,utf-8;q=0.7,*;q=0.7
Keep-Alive: 300
Connection: keep-alive
Content-Type: application/x-www-form-urlencoded
Content-Length: 37

username=test%40foo.xx&password=hello
```

We are not interested in the first part (the HTTP header, since the browser module will take care of this) but in the last part, the posted data. Since the fields of the form were username and password, the posted data is `username=test%40foo.xx&password=hello`. Now we want to reproduce the same HTTP request with our plugin. This is the simple code that will do just that.

```
-- -----
-- Must login
function pass(pstate,password)
    foo_globals.password = password

    print("*** the user inserted '"..password..
        "' as the password for '"..foo_globals.username..'")

    -- create a new browser
    local b = browser.new()

    -- store the browser object in globals
    foo_globals.browser = b

    -- create the data to post
    local post_data = string.format("username=%s&password=%s",
        foo_globals.username,foo_globals.password)
    -- the uri to post to
    local post_uri = "http://localhost:3000/"
```



```
-- post it
local file,err = nil, nil

file,err = b:post_uri(post_uri,post_data)

print("we received this webpage: ".. file)
return POPSERVER_ERR_AUTH
end
```

First we create a browser object, than we build the `post_uri` and `post_data` using a simple `string.format` (printf-like function). And this is the resulting request

```
POST / HTTP/1.1
User-Agent: Mozilla/5.0 (X11; U; Linux i686; en-US; rv:1.6) Gecko/20040322 Firefox/0.8
Pragma: no-cache
Accept: */*
Host: localhost
Content-Length: 35
Content-Type: application/x-www-form-urlencoded

username=test@foo.xx&password=hello
```

that is essentially the same (we should url-encode the post data with `curl.escape()`) we wanted to do. We saved the browser object to the global table, since we want to use the same browser all the time.

Now that we have logged in, we want to check the resulting page, and maybe extract a session ID that will be used later. This is the code to extract the session id and the HTML page we have received in response to the login request

```
... the same as before here ...


print("we received this webpage: ".. file)

-- search the session ID
local __,id = string.find(file,"session_id=(%w+)")

if id == nil then
    return POPSERVER_ERR_AUTH
end

foo_globals.session_id = id
return POPSERVER_ERR_OK
end
```

and this is the returned web page.



```
<html>
<head>
<title>foo.xx webmail</title>
</head>
<body style="background-color : grey; color : white">
<h1>Webmail - test@foo.xx</h1>
Login done! click here to view the inbox folder.
<a href="http://localhost:3000/inbox.php?session_id=ABCD1234">inbox</a>
</body>
</html>
```

Note that we extracted the session ID string using `string.find(file, "session_id=(%w+)")`. This is a really important function in the lua library and, even if it is described in the lua tutorial at [HTTP://lua-user.org](http://lua-user.org), we will talk a bit about captures here. Look at the page source. We are interested in the line

```
<a href="HTTP://localhost:3000/inbox.php?session_id=ABCD1234">inbox</a>
```

that contains the `session_id` we want to capture. Our expression is `session_id=(%w+)` that means we want to match all the strings that start with `session_id=` and then continue with one or more alphanumerical character. Since we wrote `%w+` in round brackets, we mean to capture the content of brackets (the alphanumerical part). So `string.find` will return 3 values, the first two are ignored (assigned to the dummy variable `_`) while the third is the captured string (in our case `ABCD1234`). The LUA tutorial at lua-users is quite good and at [HTTP://sf.net/projects/lua-users](http://sf.net/projects/lua-users) you can find the LUA short reference that is a summary of all standard lua functions and is a really good piece of paper (so many thanks to Enrico Colombini). If you really like LUA you should buy THE book about LUA called "*Programming in Lua*" by Roberto Ierusalimsky (consider it the K&R for LUA).

7.2.3 (step 3) getting the list of messages

Now we have to implement the `stat()` function. The `stat` is probably the most important function. It must retrieve the list of messages in the webmail and their UIDL and size. In our example we will use the `mlex` module to grab the

important info from the page, but you can use the string LUA module to do the same with captures. This is our inbox page



and this is the HTML body (only the first 2 messages are reported)

```
<h1>test@foo.xx - inbox (1/2)</h1>
<form name="inbox" method="post" action="/delete.php">
<input type="hidden" name="session_id" value="ABCD1234">
<table>
<tr><th>From</th><th>subject</th><th>size</th><th>date</th></tr>
<tr>
  <td><b>friend1@foo1.xx</b></td>
  <td><b><a href="/read.php?session_id=ABCD1234&uidl=123">ok!</a></b></td>
  <td><b>20KB</b></td>
  <td><b>today</b></td>
  <td><input type="checkbox" name="check_123"></td>
</tr>
<tr>
  <td>friend2@foo2.xx</td>
  <td><a href="/read.php?session_id=ABCD1234&uidl=124">Re: hi!</a></td>
  <td>12KB</td>
  <td>yesterday</td>
  <td><input type="checkbox" name="check_124"></td>
</tr>
</table>
<input type="submit" value="delete marked">
</form>
<a href="/inbox.php?session_id=ABCD1234&page=2">go to next page</a>
</body>
```

We have retrieved the HTML using the browser and the `get_uri()` method (remember the uri for the inbox was in the login page). As you can see the

messages are in a table and this table has the same structure for each message. This is the place in which you may use `mlex`. Just take all the stuff between `<tr>` and `</tr>` of a message row and delete all but the tags name. Then replace every empty space (we call space the string between two tags) with a `".*"`. This is what we have obtained (it should be all in the same line, here is wrapped for lack of space) from the first message.

```
.*<tr>.*<td>.*<b>.*</b>.*</td>.*<td>.*<b>.*<a>.*</a>.*</b>.*</td>.*
<td>.*<b>.*</b>.*</td>.*<td>.*<b>.*</b>.*</td>.*
<td>.*<input>.*</td>.*</tr>
```

This expression is used to match the table row containing info about the message. Now cut and paste the line and replace every space and every tag with `O` (the letter, not the digit `0`) or `X`. Put an `X` in the interesting fields (in our example the size and the input tag, that contains the message uidl).

```
0<O>0<O>0<O>0<O>0<O>0<O>0<O>0<O>0<O>0<O>0<O>0<O>0
<O>0<O>X<O>0<O>0<O>0<O>0<O>0<O>0<O>0
<O>0<X>0<O>0<O>0
```

While the first expression will be used to match the table row, this one will be used to extract the important fields. This is the code that starts `mlex` on the HTML and fills the `popstate` data structure with the captured data.

```
-----
-- Fill the number of messages and their size
function stat(pstate)
    local file,err = nil, nil
    local b = foo_globals.browser
    file,err = b:get_uri("http://localhost:3000/inbox.php?session_id"..
        foo_globals.session_id)
    local e = ".*<tr>.*<td>.*<b>.*</b>.*</td>.*<td>.*<b>.*<a>"..
        ".*</a>.*</b>.*</td>.*<td>.*<b>.*</b>.*</td>.*<td>.*".."
        "<b>.*</b>.*</td>.*<td>.*<input>.*</td>.*</tr>"
    local g = "0<O>0<O>0<O>0<O>0<O>0<O>0<O>0<O>0<O>0<O>0<O>0<O>0"
        "<O>0<O>X<O>0<O>0<O>0<O>0<O>0<O>0<O>0<O>0<X>0<O>0<O>0"
    local x = mlex.match(file,e,g)
    --debug print
    x:print()

    set_popstate_nummesg(pstate,x:count())
    for i=1,x:count() do
        local __,_,size = string.find(x:get(0,i-1),("(%d+)")
        local __,_,size_mult_k = string.find(x:get(0,i-1),("([Kk][Bb])"))
        local __,_,uidl = string.find(x:get(1,i-1),"check_(%d+)")
    end
end
```

```

        if size_mult_k ~= nil then
            size = size * 1024
        end
        if size_mult_m ~= nil then
            size = size * 1024 * 1024
        end

        set_mailmessage_size(pstate,i,size)
        set_mailmessage_uidl(pstate,i,uidl)
    end

    return POPSERVER_ERR_OK
end

```

The result of `x:print()` is this one

```
{'20KB','input type="checkbox" name="check_123"'} }
```

and the telnet session follows

```

+OK FreePOPs/0.0.11 pop3 server ready
user test@foo.xx
+OK PLEASE ENTER PASSWORD
pass secret
+OK ACCESS ALLOWED
stat
+OK 1 20480
quit
+OK BYE BYE, UPDATING

```

We have not listed here how we added the dummy `return POPSERVER_ERR_OK` line to the `quit()` function. The source code listed before uses `mlex` to extract the two interesting strings, then parses them searching for the size and the size multiplier and the uidl. Then sets the mail message attributes. But here you can see that we just matched the first message. To match the other messages we have to inform the `mlex` module that the `` tag is optional (you can see that only the first message is in bold). so we change the expressions to

```

.*<tr>.*<td>[.]*{b}.*{/b}[.]*</td>.*<td>[.]*{b}.*<a>.*</a>.*{/b}[.]*</td>.*
<td>[.]*{b}.*{/b}[.]*</td>.*<td>[.]*{b}.*{/b}[.]*</td>.*
<td>.*<input>.*</td>.*</tr>

```

and

```

0<0>0<0>[0]{0}0{0}[0]<0>0<0>[0]{0}0<0>0<0>0{0}[0]<0>0
<0>[0]{0}x{0}[0]<0>0<0>[0]{0}0{0}[0]<0>0
<0>0<x>0<0>0<0>

```

Now the stat command responds with +OK 4 45056 and the debug print is

```

{'20KB','input type="checkbox" name="check_123"'}
{'12KB','input type="checkbox" name="check_124"'}
{'10KB','input type="checkbox" name="check_125"'}
{'2KB','input type="checkbox" name="check_126"'}

```

Now we have a proper function stat that fill the popstate data structure with the info the popserver needs to respond to a stat request. Since the list, uidl, list_all and uidl_all requests can be satisfied with the same data we will use the standard function provided by the common.lua module. It will be explained in the next step, but we have to add 2 important lines to the stat function. They avoid a double call to the stat() function.

```

function stat(pstate)
    if foo_globals.stat_done == true then return POPSERVER_ERR_OK end

    ... the same code here ...

    foo_globals.stat_done = true
    return POPSERVER_ERR_OK
end

```

The most important function is done, but a lot of notes must be written here. First, mlex is really comfortable sometimes, but you may find more helpful using the lua string library or the regex library (posix extended regular expressions) to reach the same point. Second, this implementation stops at the first inbox page. You should visit all the inbox pages maybe using the do_until() function in the support.lua library (that will be briefly described at the end of this tutorial). Third we make no error checking. For example the file variable may be nil and we must check these things to make a good plugin.

7.2.4 (step 4) the common functions

The common module gives us some pre-cooked functions that depend only on a well implemented stat() (I mean a stat than can be called more than once). This is our implementation of these functions

```

-----
-- Fill msg uidl field
function uidl(pstate,msg) return common.uidl(pstate,msg) end

-----
-- Fill all messages uidl field
function uidl_all(pstate) return common.uidl_all(pstate) end

-----
-- Fill msg size
function list(pstate,msg) return common.list(pstate,msg) end

-----
-- Fill all messages size
function list_all(pstate) return common.list_all(pstate) end

-----
-- Unflag each message marked for deletion
function rset(pstate) return common.rset(pstate) end

-----
-- Mark msg for deletion
function dele(pstate,msg) return common.dele(pstate,msg) end

-----
-- Do nothing
function noop(pstate) return common.noop(pstate) end

```

but first add the common module loading code to your `init()` function

```

... the same code ..

-- the common module
if freepops.dofile("common.lua") == nil then
    return POPSERVER_ERR_UNKNOWN
end

-- checks on globals
freepops.set_sanity_checks()

return POPSERVER_ERR_OK
end

```

7.2.5 (step 5) deleting messages

Deleting messages is usually a normal post and an example of the `post_data` is `session_id=ABCD1234&check_124=on&check_126=on`. The code follows

```

-----
-- Update the mailbox status and quit
function quit_update(pstate)
    -- we need the stat
    local st = stat(pstate)
    if st ~= POPSERVER_ERR_OK then return st end

    -- shorten names, not really important
    local b = foo_globals.b
    local post_uri = b:wherearewe() .. "/delete.php"
    local session_id = foo_globals.session_id
    local post_data = "session_id=" .. session_id .. "&"

    -- here we need the stat, we build the uri and we check if we
    -- need to delete something

    local delete_something = false;
    for i=1,get_popstate_nummsg(pstate) do
        if get_mailmessage_flag(pstate,i,MAILMESSAGE_DELETE) then
            get_mailmessage_uidl(pstate,i).. "=on&"
            delete_something = true
        end
    end

    if delete_something then
        b:post_uri(post_uri,post_data)
    end

    return POPSERVER_ERR_OK
end

```

Consider we do the post only if at least one message is marked for deletion. Another important think to keep in mind is that making only one post for all messages is better than making a single post for each message. When it is possible you should reduce the number of HTTP requests as much as you can since it is here we move FreePOPs from a rabbit to a tortoise.

7.2.6 (step 6) downloading messages

You may ask why I talk about this only at point 6, while having the mail is probably what you want from a plugin. Implementing the `retr()` function is usually simple. It really depends on the webmail, but here we will talk of the simple case, while at the end of the tutorial you will see how to deal with complex webmails. The simple case is the one in which the webmail has a save message button. And the saved message is a plain text file containing both the

header and the body of the message. There are only two interesting points in this case, first big messages, last the dot issue.

Big messages are a cause of timeout. Yes, the most simple way of downloading a message is calling `b:get_uri()` and store the message in a variable, and then send it to the mail client with `popserver_callback()`. But think of a 5MB mail, downloaded with a 640Kb DSL connection, at full 80KB speed, it takes 64 seconds to download. This means your plugin will not send any data to the mail client for more than one minute. And this will make the mail client to disconnect from FreePOPs thinking the POP3 server is dead. So we must send the data to the mail client as soon as we can. For this we have the `b:pipe_uri()` function that calls a callback whenever it has some fresh data. The following code is the callback factory function, that creates a new callback to pass to the `pipe_uri` browser method.

```
-----
-- The callback factory for retr
--
function retr_cb(data)
    local a = stringhack.new()
    return function(s,len)
        s = a:dothack(s).."\0"
        popserver_callback(s,data)
        return len,nil
    end
end
```

Here you see that the callback simply uses `popserver_callback()` to pass the data to the mail client, but before doing this it mangles the data with the `stringhack`. But this is the second interesting point.

The POP3 protocol should end the retr command answer with a line that contains only 3 bytes, `“.\r\n”`. But what if a line, inside the mail body, is a simple point? We have to escape it to `“..\r\n”`. This is not so hard, a `string.gsub(s,“\r\n.\r\n”,“\r\n..\r\n”)` is all we need... but not in the case of callbacks. The send callback will be called with some fresh data, and called more than one time if the mail is big. And if the searched pattern is truncated between two calls the `string.gsub()` method will fail. This is why the `stringhack` module helps us. The `a` object lives as long as the callback function will be called (see the closure page of the lua tutorial) and will keep in mind that the searched pattern may be truncated.

Finally the `retr()` code

```
-----
-- Get message msg, must call
-- popserver_callback to send the data
```

```

function retr(pstate,msg,pdata)
    -- we need the stat
    local st = stat(pstate)
    if st ~= POPSERVER_ERR_OK then return st end

    -- the callback
    local cb = retr_cb(data)

    -- some local stuff
    local session_id = foo_globals.session_id
    local b = internal_state.b
    local uri = b:wherearewe() .. "/download.php?session_id"..session_id..
        "&message="..get_mailmessage_uidl(pstate,msg)

    -- tell the browser to pipe the uri using cb
    local f,rc = b:pipe_uri(uri,cb)
    if not f then
        log.error_print("Asking for "..uri.."\\n")
        log.error_print(rc.."\\n")
        return POPSERVER_ERR_NETWORK
    end
end
end

```

7.2.7 (step 7) test it

Making a good plugin needs a lot of testing. You should ask for beta testers at the FreePOPs forum ([HTTP://freepops.diludovico.it](http://freepops.diludovico.it)) and ask the software authors to include it in the main distribution. You should also read the webmail contract, check if there is something like “*I’ll never use webmail->pop3 server to read my mail*” and send a copy to the authors of the software.

7.2.8 (step 8) the so mentioned last part of the tutorial

There are a lot of things we have omitted here.

The multi-page stat is the real good implementation for `stat()`. We mentioned before that our implementation lists only the messages in the first page. The code for parsing and extracting interesting info from a page is already written, we simply need a function that checks if we are in the last page and if not it changes the value of a `uri` variable. The `uri` variable will be used by the `fetch` function. In this case you should use the `support` module with the `do_until` cycle. This is a simple example of `do_until()`

```

-- -----
-- Fill the number of messages and their size

```

```

function stat(pstate)
    ... some code as before ...

    -- this string will contain the uri to get. it may be updated by
    -- the check_f function, see later
    local uri = string.format(libero_string.first,popserver,session_id)

    -- The action for do_until
    --
    -- uses mlex to extract all the messages uidl and size
    local function action_f (s)
        -- calls match on the page s, with the mlexpressions
        -- state and statG
        local x = mlex.match(s,e,g)

        -- the number of results
        local n = x:count()

        if n == 0 then return true,nil end

        -- this is not really needed since the structure
        -- grows automatically... maybe... don't remember now
        local nmesg_old = get_popstate_nummesg(pstate)
        local nmesg = nmesg_old + n
        set_popstate_nummesg(pstate,nmesg)

        -- gets all the results and puts them in the popstate structure
        ... some code as before ...

        set_mailmessage_size(pstate,i+nmesg_old,size)
        set_mailmessage_uidl(pstate,i+nmesg_old,uidl)
    end

    return true,nil
end

-- check must control if we are not in the last page and
-- eventually change uri to tell retrieve_f the next page to retrieve
local function check_f (s)
    local tmp1,tmp2 = string.find(s,next_check)
    if tmp1 ~= nil then
        -- change retrieve behaviour
        uri = "--build the uri for the next page--"

        -- continue the loop
        return false
    else

```

```

        return true
    end
end

-- this is simple and uri-dependent
local function retrieve_f ()
    local f,err = b:get_uri(uri)
    if f == nil then
        return f,err
    end

    local __,_,c = string.find(f,"--timeout string--")
    if c ~= nil then
        internal_state.login_done = nil
        session.remove(key())
        local rc = libero_login()
        if rc ~= POPSERVER_ERR_OK then
            return nil,"Session ended,unable to recover"

            uri = "--uri for the first page--"
            return b:get_uri(uri)
        end
    end

    return f,err
end

-- initialize the data structure
set_popstate_nummesg(pstate,0)

-- do it
if not support.do_until(retrieve_f,check_f,action_f) then
    log.error_print("Stat failed\n")
    session.remove(key())
    return POPSERVER_ERR_UNKNOWN
end

-- save the computed values
internal_state["stat_done"] = true
return POPSERVER_ERR_OK
end

```

The only strange things are the retrieve function and the session saving stuff. Since webmail sometimes timeout you should check if the retrieved page is valid or not, and eventually retry the login. The session saving is the next issue.

Saving the session is the way to make FreePOPs really similar to a browser.

This means the next time you check the mail FreePOPs will simply reload the inbox page and doesn't login again. To do this you need a `key()` function that gives a unique id for each session

```
-----  
-- The key used to store session info  
--  
-- This key must be unique for all webmails, since the session pool is one  
-- for all the webmails  
--  
function key()  
    return foo_globals.username .. foo_globals.password  
end
```

and a `foo_globals` serialization function

```
-----  
-- Serialize the internal state  
--  
-- serial.serialize is not enough powerful to correctly serialize the  
-- internal state. The field b is the problem. b is an object. This means  
-- that it is a table (and no problem for this) that has some field that are  
-- pointers to functions. this is the problem. there is no easy way for the  
-- serial module to know how to serialize this. so we call b:serialize  
-- method by hand hacking a bit on names  
--  
function serialize_state()  
    internal_state.stat_done = false;  
    return serial.serialize("foo_globals",foo_globals) ..  
        internal_state.b:serialize("foo_globals.b")  
end
```

Now you have to tell FreePOPs to save the state in the `quit_update()` function and load it back in the `pass()` one. This is the new `pass()` structure

```
function pass(pstate,password)  
    -- save the password  
    internal_state.password = password  
  
    -- eventually load session  
    local s = session.load_lock(key())  
  
    -- check if loaded properly
```

```

    if s ~= nil then
        -- "\a" means locked
        if s == "\a" then
            log.say("Session for "..internal_state.name..
                " is already locked\n")
            return POPSERVER_ERR_LOCKED
        end

        -- load the session
        local c,err = loadstring(s)
        if not c then
            log.error_print("Unable to load saved session: "..err)
            return foo_login()
        end

        -- exec the code loaded from the session string
        c()

        log.say("Session loaded for " .. internal_state.name .. "@" ..
            internal_state.domain ..
            "(" .. internal_state.session_id .. ")\n")

        return POPSERVER_ERR_OK
    else
        -- call the login procedure
        return foo_login()
    end
end
end

```

where `foo_login()` is the old `pass()` function with minor changes. Don't forget to call `session.unlock(key())` in the `quit()` function, since you have to release the session in case of failure (and `quit()` is called here) and to save the session in `quit_update()`

```

-- save fails if it is already saved
session.save(key(),serialize_state(),session.OVERWRITE)
-- unlock is useless if it have just been saved, but if we save
-- without overwriting the session must be unlocked manually
-- since it would fail instead overwriting
session.unlock(key())

```

The `top()` function is a complex thing. I'll not describe it in a complete way, but I suggest you to look at the `libero.lua` plugin if the web server that sends you the message source supports the "Range:" HTTP request field, or the `tin.lua` plugin if the server needs to be interrupted in a bad way.

Remember that the `top()` need someone that counts the lines and here we have again the `stringhack` module that counts and eventually purges some lines.

The javascript is the hell of webmails. Javascript can do anything and you have to read them to emulate what they do. For example they may add some cookies (and you'll have to do this by hand with the `b:add_cookie()` as in `tin.lua`) or they may change some form fields (like in the `libero.lua` login load balancing code).

The cookies are sweet enough for us, since the browser module will handle them for us.

The standard files are really system dependent. Under Windows you'll have to constantly look at the `stderr.txt` and `stdout.txt`, while under Unix you will just have to start it with the `-w` switch and look at the console.

The brute force is called `ethereal`. Sometimes things don't work in the right way and the only way to debug them is to activate curl debugging to see what FreePOPs does (`b:curl:setopt(curl.OPT_VERBOSE,1)`) and sniff what a real browser does.

The open source way is the best way of having a good quality peace of software. This means you'll have to release really often your plugin in the development phase and interact much with your testers. Trust me it works, or read "*The cathedral and the bazaar*" by Eric Raymond.

The mimer module is really beta at the time of this tutorial, but is what you need if you are in the unlucky case of a webmail that has no save message button. The `lycos.lua` plugin is an example of what it can do. The main interesting function is `mimer.pipe_msg()` that takes a message header, a text body (in html o plain text format) and some attachments uris, that are downloaded on the fly, composed into a proper mail message and piped to the mail client.

8 Submitting a bug

When you have problems or you think you have found a bug, please follow strictly this *iter*:

1. Update to the most recent version of FreePOPs.
2. Try to reproduce the bug, if the bug is not easily reproducible we are out of luck. Something can still be tried: if the software crashed you could compile it from the sources, install `valgrind`, run `freepopsd` with `valgrind` and hope the error messages are interesting.

3. Clean the log files
4. Start FreePOPs with the -w switch
5. Reproduce the bug
6. Send to the developers the log, plus any other info like your system type and how to reproduce this bug.

9 Authors

This manual has been written by Enrico Tassi <gareuselesinge@users.sourceforge.net> and revised and translated by Nicola Cocchiaro <ncocchiaro@users.sourceforge.net>

9.1 Developers

FreePOPs is developed by:

- Enrico Tassi <gareuselesinge@users.sourceforge.net>
- Alessio Caprari <alessiofender@users.sourceforge.net>
- Nicola Cocchiaro <ncocchiaro@users.sourceforge.net>
- Simone Vellei <simone_vellei@users.sourceforge.net>

LiberoPOPs is developed by:

- Enrico Tassi <gareuselesinge@users.sourceforge.net>
- Alessio Caprari <alessiofender@users.sourceforge.net>
- Nicola Cocchiaro <ncocchiaro@users.sourceforge.net>
- Simone Vellei <simone_vellei@users.sourceforge.net>
- Giacomo Tenaglia <sonicsmith@users.sourceforge.net>

10 Thanks

Special thanks goes to the users who tested the software, to the hackers who made it possible to have a free and reliable development environment as the Debian GNU/Linux system.